

Agentic AI for Healthcare Content Management: From Policy Interpretation to Auditable Rule Generation

Yue Shen

Topic: Content Management in Health Care

Healthcare content management includes the lifecycle of billing and coding policies, clinical guidelines, payer-provider contracts, and related operational rules. The central challenge is that this content is high-volume, frequently changing, and written in natural language, while payer operations require precise, machine-executable logic. This creates a near-term opportunity for agentic AI systems that can interpret policy text, draft structured rules, and support auditable human review. This report defines the concept, analyzes current trends, weighs the associated opportunities and threats, and proposes a strategic pilot for Cotiviti to consider.

Defining the Concept

In this context, content management is not just document storage. It is the operational process of transforming healthcare policy content into reviewable, testable, and auditable business logic. Payers apply that logic to adjudicate claims correctly, the discipline known as payment integrity. Today, much of that translation remains manual: analysts read coverage and coding policy and hand-convert it into rules, a process that is slow, costly, and difficult to keep current as regulations change. The stakes are considerable. CMS reports an estimated 31.70 billion dollars in improper Medicare fee-for-service payments in fiscal year 2024, with insufficient documentation the leading contributing cause (Centers for Medicare & Medicaid Services, 2024). Accurate, current, and machine-executable rules are therefore central to both cost control and regulatory compliance.

Trends

Generative AI, and increasingly agentic AI, begins to automate the pipeline from policy interpretation to rule generation. Retrieval-augmented generation (RAG) grounds a language model in an external knowledge source so that it produces cited, verifiable output rather than unsupported text. A 2025 systematic review of RAG in healthcare finds that grounding reduces hallucination and improves the transparency of sources, although standardized evaluation frameworks remain immature (Amugongo et al., 2025). Structured extraction through function calling converts policy prose into machine-readable rules, an approach often described as policy as code. Agentic systems extend this from question answering toward workflow automation: rather than only answering a query, the system retrieves evidence, drafts rules, tests them against sample cases, and flags exceptions. The same pattern supports change management, where a model compares policy versions and surfaces downstream impact. (A visual comparison of the traditional, question-answering, and agentic approaches appears in the accompanying presentation.) Industry momentum is real. Cotiviti identifies AI as a key investment area for payment integrity and partners with the Responsible AI Institute to govern its deployment (Responsible AI Institute, 2025).

Opportunities

The projected benefits center on speed, accuracy, and auditability. Automating a first draft of rule logic can compress rule development and update cycles, an estimated potential improvement as regulations change. More consistent and complete rule coverage can reduce improper payments, an estimated potential gain given the tens of billions of dollars at stake each year. Analysts shift from building rules to reviewing AI-drafted candidates, which can raise throughput while preserving expert judgment. Citations create an audit trail from each decision back to the source policy, which is essential in a regulated environment. The same architecture extends

naturally to policy change management and knowledge assistance. Industry analysis suggests that AI could automate a projected 8 to 30 percent of health plan administrative costs and reduce overall medical costs by a projected 0.4 to 1.7 percent (Cotiviti, 2024).

Threats

The risks are real, and they concentrate on accuracy and governance. A misextracted rule in a compliance-critical setting can produce an incorrect denial or payment, so extraction quality requires measurement rather than assumption. Allowing a model to decide claims directly is hazardous; final rule application should remain deterministic and open to human review. Governance, accountability, and explainability matter, because regulators and members expect transparency. Data governance and intellectual property require care, since proprietary specifications and protected health information cannot be handled casually. Model drift is a standing concern, because policies change and extraction pipelines require ongoing revalidation. Adoption also remains uneven across payers, with integration and security cited as leading barriers (Cotiviti, 2025). The central question is therefore not whether to adopt agentic AI, but how to adopt it responsibly.

Strategic Recommendation

Cotiviti should pilot a human-in-the-loop policy-to-rule copilot for payment accuracy analysts. As a payment integrity partner to 23 of the top 25 national payers (Cotiviti, 2026), Cotiviti is well positioned to test this capability at meaningful scale.

- **What it does:** the copilot ingests coverage or coding policy, retrieves relevant passages, drafts structured rules with citations, tests those rules against sample claims, and routes ambiguous cases to analysts.
- **Why this design:** language models are useful for interpretation and extraction, while deterministic engines handle final rule application to preserve auditability and consistency. This mirrors Cotiviti's own stated principle that AI augments rather than replaces expert judgment (Responsible AI Institute, 2025).
- **How to phase it:** the pilot starts with a narrow, high-volume policy area; measures extraction accuracy, analyst acceptance rate, cycle-time reduction, and exception-review quality; and then expands to policy change management and broader content workflows.

This recommendation is near-term and governable rather than a moonshot. It aligns with Cotiviti's payment accuracy mission, its quality and Stars focus, and its responsible AI commitments, and its value is a projected reduction in cycle time and improper payments supported by a defensible audit trail. A working prototype that demonstrates this design, including retrieval, cited rule extraction, deterministic adjudication, and human-review routing, accompanies this report.

References

- Amugongo, L. M., Mascheroni, P., Brooks, S., Doering, S., & Seidel, J. (2025). Retrieval augmented generation for large language models in healthcare: A systematic review. *PLOS Digital Health*, 4(6), e0000877. <https://doi.org/10.1371/journal.pdig.0000877>
- Centers for Medicare & Medicaid Services. (2024). *Fiscal year 2024 improper payments fact sheet*. <https://www.cms.gov/newsroom/fact-sheets/fiscal-year-2024-improper-payments-fact-sheet>
- Cotiviti. (2024). *The evolving role of artificial intelligence in payment integrity* [eBook]. <https://resources.cotiviti.com/payment-integrity/the-evolving-role-of-artificial-intelligence-in-payment-integrity>
- Cotiviti. (2025). *The practical landscape of AI in payment integrity and VBP*. <https://resources.cotiviti.com/payment-integrity/the-practical-landscape-of-ai-in-payment-integrity-and-vbp>
- Cotiviti. (2026, July). *Cotiviti recognized by Everest Group as the highest-designated leader for market impact in the 2026 pre-payment integrity solutions PEAK Matrix assessment* [Press release]. <https://www.businesswire.com/news/home/20260701943009/en/>
- Responsible AI Institute. (2025). *Cotiviti joins Responsible AI Institute as a new member*. <https://www.responsible.ai/news/cotiviti-joins-responsible-ai-institute-as-a-new-member/>